

# Business Insights and Recommendations Report: Loan Approval Prediction

**Prepared by:** Edwin Sunday Umana

**Role:** Data Science Intern, AnalystLab Africa

**Date:** 20 August 2026 (Week 3)

**Batch:** D

## 1. Executive Summary

The objective of this analysis was to investigate the historical loan application dataset to uncover the primary drivers of loan approval and prepare a robust dataset for predictive machine learning. Through advanced exploratory data analysis, rigorous statistical hypothesis testing, and domain-driven feature engineering, I identified that an applicant's past credit history and their relative financial burden (loan-to-income ratio) are the most critical determinants of loan approval. Raw income alone is a surprisingly weak predictor. By engineering targeted risk flags and eliminating redundant variables, I have refined the dataset into a high-quality foundation for automated credit risk modelling in Week 4.

## 2. Key Analytical Findings

My investigation yielded several significant patterns and statistically validated insights:

**Credit History is the Ultimate Gatekeeper:** Through a Chi-Square Test of Independence, I mathematically proved that past credit behaviour is heavily dependent on loan approval status. Applicants with a bad credit history (0) face an overwhelmingly high rejection rate, making it the single strongest positive correlate in the dataset.

**Financial Burden Outweighs Raw Wealth:** Using the Mann-Whitney U Test, I confirmed that the `Loan_to_Income_Ratio` is significantly higher for rejected loans. Furthermore, my visual analysis revealed that while high-income applicants exist in the rejected pool, they were often rejected because they requested disproportionately massive loan amounts.

**Geographic Income Disparities:** A One-Way ANOVA test confirmed that applicant incomes vary significantly across Rural, Semiurban, and Urban property areas. However, property area alone does not dictate approval; it merely segments the applicant pool's financial baseline.

**The Power of Engineered Features:** I discovered that original features like `ApplicantIncome` and `LoanAmount` suffer from multicollinearity and fail to capture the full risk profile. By engineering `TotalIncome`, `Income_Per_Dependent`, and the `High_Risk_Flag`, I created much stronger, normalized signals that clearly separate approved from rejected classes.

### 3. Business Implications

For a financial institution, these findings highlight the danger of relying on manual underwriting or simplistic income thresholds.

Approving loans based purely on a high salary is risky if the applicant's debt-to-income ratio is unsustainable.

The heavy reliance on `Credit_History` implies that the bank's current risk models are highly conservative regarding past defaults.

The success of the `High_Risk_Flag` (which combines bad credit and high financial burden) demonstrates that compound risk indicators are far more effective at flagging potential defaulters than looking at variables in isolation.

### 4. Evidence-Based Recommendations

Based on my analysis, I recommend the following strategic actions for the lending institution:

**Implement Automated Fast-Tracking:** Applications with a perfect `Credit_History` (1) and a standard `High_Risk_Flag` (0) should be routed to an automated instant-approval system. This will drastically reduce processing time and improve customer satisfaction for low-risk applicants.

**Enforce a Strict Loan-to-Income Threshold:** The bank should establish a hard cap on the `Loan_to_Income_Ratio`. Any application exceeding this threshold should trigger an automatic rejection or require manual underwriting with additional collateral, regardless of the applicant's raw income.

**Tailor Products by Geography:** Since ANOVA results proved income baselines differ by `Property_Area`, the marketing and product teams should design region-specific loan products (e.g., micro-loans for specific rural demographics vs. premium mortgages for urban centres).

**Adopt the Refined Feature Set for AI:** The machine learning models developed in Week 4 should rely heavily on the engineered interaction features (`High_Risk_Flag`, `EMI`, `Loan_to_Income_Ratio`) rather than raw numerical columns, as they provide a much cleaner, noise-free signal for algorithms like Logistic Regression and Random Forests.

### 5. Dataset Limitations

While my analysis provides strong directional insights, I must acknowledge a few limitations in the current dataset:

**Lack of Granular Credit Scores:** The `Credit_History` variable is binary (0 or 1). A real-world credit score (e.g., FICO score ranging from 300-850) would allow for much more nuanced risk tiering.

**Sample Size:** The training dataset contains only 614 records. While sufficient for initial statistical testing, a larger dataset would be required to train deep learning models or capture rare edge-case defaults.

**Missing Interest Rates:** The dataset does not include the interest rate assigned to the loan, which is a critical factor in calculating the true monthly EMI burden and the bank's actual profitability.

### 6. Next Steps

The immediate next step is to transition from statistical analysis to predictive modelling (Week 4). I will use the `final_modelling_dataset.csv`—which has been scaled, encoded, and stripped of redundant

multicollinear features—to train and evaluate multiple classification algorithms (such as Logistic Regression, Random Forest, and XGBoost). The goal will be to deploy a model that can accurately predict the `Loan_Status` of unseen test applications, thereby automating the underwriting process.